



CS698R: Deep Reinforcement Learning Project

FORAGING IN THE FIELD

Mentors: Adithya, Ajita

Group-5:

Tabish Ahmad - 21111060

Prakhar Srivastava - 170486

Areeb Ahmad -180135

Mohd Niyas P -170394

Samarth Mehrotra -20128408

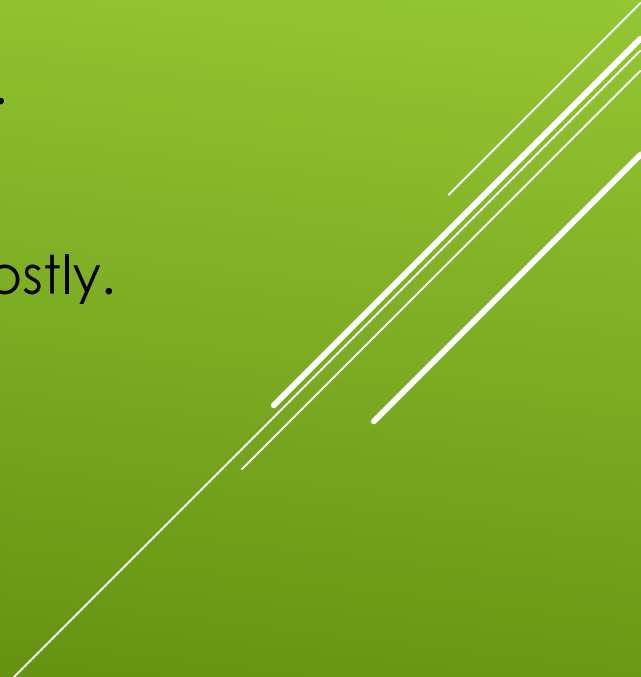
PROBLEM DESCRIPTION:

- ▶ There is a field with certain number of patches, and each patch has certain number of berries of various sizes.
- ▶ Reward is proportional to the size of the berry.
- ▶ There is a pre-defined variable draining rate proportional to distance travelled which depletes the total reward accumulated.
- ▶ Continuous movement or having the option to stay?
- ▶ Field constant or changing ?

AIM:

- ▶ Maximize the reward remaining at the end?
- ▶ Maximize the total cumulative reward collected during 5 mins?
- ▶ Both?

MOTIVATION:

- ▶ Although most decision research concerns choice between simultaneously presented options, in many situations options are encountered sequentially.
 - ▶ The decision is whether to exploit an option or explore for a better one.
 - ▶ All animals must forage food for survival but doing so is energetically costly.
 - ▶ But we know little about process involved in decision making.
- 

LITERATURE SURVEY:

- ▶ **Information foraging** : - Application of optimal foraging theory to understand human behaviour while searching for information on internet.
- ▶ **Marginal Value Theorem** : It describes the behaviour of an optimally foraging individual in an environment where resources (berries) are located in discrete patches separated by areas with no resources.
- ▶ **Observation of Animals** : - Species like great tits and screaming hairy armadillos have been studied and their foraging patterns have been tested for optimality.
- ▶ **NEAT**: Neural Evolution of Augmented Topologies. A genetic algorithm which evolves the neural networks over generations.

ENVIRONMENT DETAILS: (RECAP)

Overall

- Field - 20,000 x 20,000 pixels
- Screen - 1920 x 1080
- Time - 5 minutes

Patch

- Size - 2600 x 2600
- Number - 10
- Patches completely inside field
- Interpatch distance - ≥ 5000
- Berry sizes - 10, 20, 30, 40
- Berry number - 20 each/patch

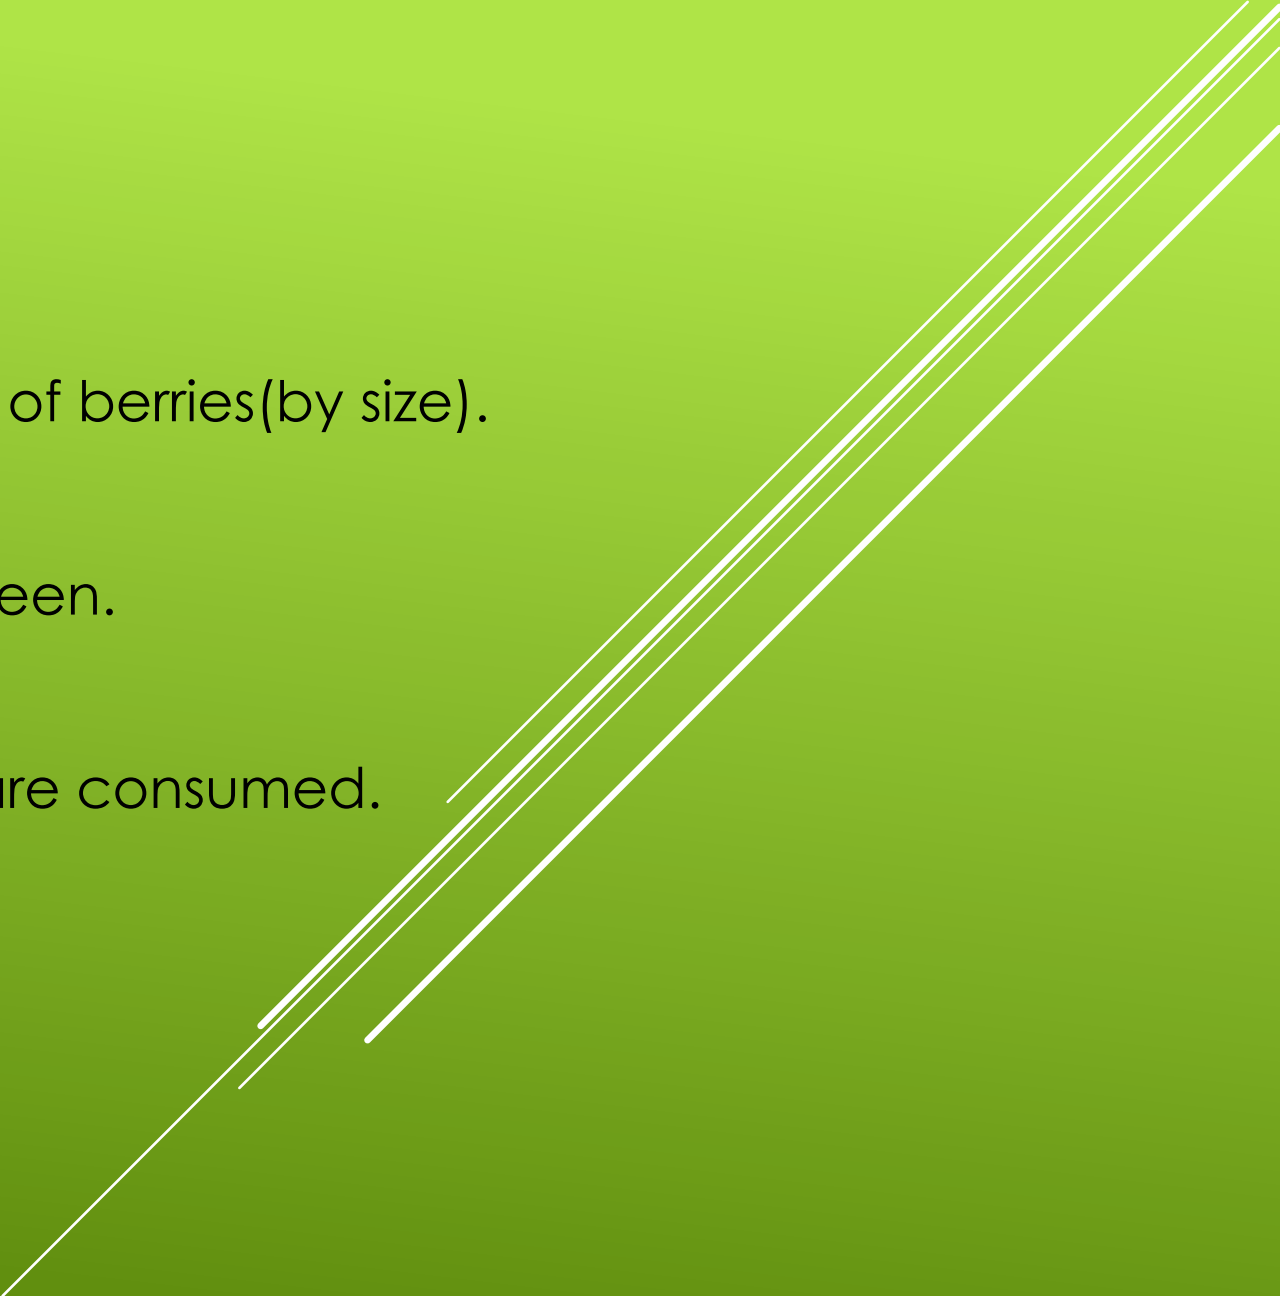
Reward

- Initial = 0.5 (out of 1)
- Drain rate = $0.5d / (120 \times 400)$ (d distance in pixels)
- Berry reward = diameter/10,000

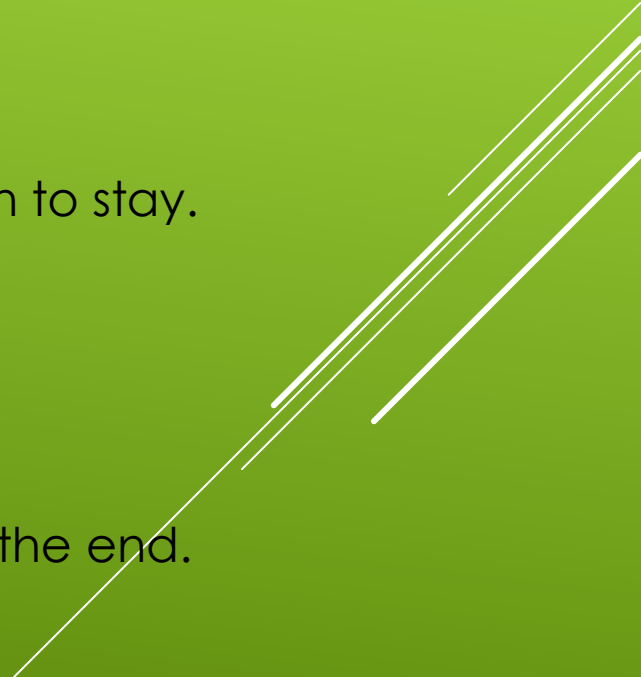
Player

- Size = 10
- Speed - 400 pixels/seconds
- Arrow key movement (8 directions)

ASSUMPTIONS

- ✓ Agent knows that there are four types of berries(by size).
 - ✓ Agent's vision is restricted to visible screen.
 - ✓ The berries don't replenish after they are consumed.
 - ✓ Patches has same densities.
- 
- Several thin, white, parallel diagonal lines are drawn across the bottom right portion of the slide, extending from the bottom edge towards the top right corner.

IMPLEMENTATION:

- Environment implemented with OpenAi gym integrated with PyGame for rendering.
 - Environment has the option human and computer(agent) play.
 - Agent is trained by the algorithm Double DQN.
 - Agent is provided two options: Continuous Action or Have an additional option to stay.
 - Continuous action helps in maximizing the cumulative reward.
 - Non-continuous action allows the agent to maximize the reward remaining at the end.
- 

Agents

```
graph TD; Agents[Agents] --> Continuous[Continuous Movement]; Agents --> NonContinuous[Non-Continuous Movement]; Agents --> Random[Random Environment];
```

Continuous
Movement

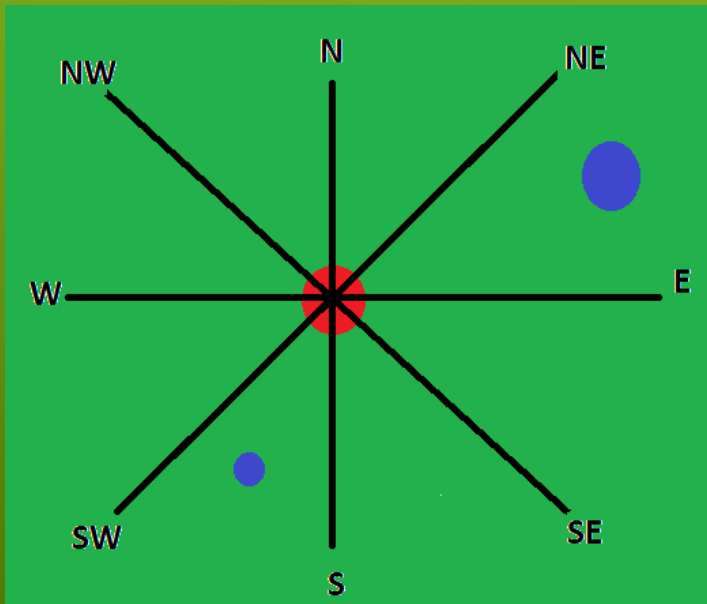
Non-
Continuous
Movement

Random
Environment

AGENT:

- State-Size = 35, Action-Size = 8 or 9 (depending on continuous or non-continuous movement).
- The input vector to our neural network is a 35 length vector. The output is a vector of size 8 or 9.
- Since we can't use CNN because of **memory issues** and **FPS** drops we took the information from the image as features and provided it to the network.
- The features are:

Consider the vector: [N, S, E, W, NE, NW, SE, SW]



- The agent scans in the 8 contained directions as far as it can see or until it sees a berry .
- The vector is filled as: [0, 0, size1/distance, 0, size1/distance, 0, 0, 0]
- Another example: [0, size2/distance, 0, 0, 0, 0, 0, size2/distance]
- There is a different vector for each size of berry. Therefore we have $8 \times 4 = 32$ states.
- **Intuition:** Is it worth travelling the distance for the corresponding size?

- 33rd state is the density. Meaning how much berries are occupying the screen at the current moment.
- 34th state is the current health remaining.
- 35th state is the current time remaining.
- **Intuition:**
 - Density helps in determining if it is a good time to leave the patch.
 - Current health and time helps in determining whether to stay in the already depleted patch instead of exploring.
 - Precisely at any point of time the agent should have the health to explore a new patch and if it doesn't the current patch should be able to provide that much resources.
 - Features are scaled so that any feature don't get too big or too small.

- Noise is added to the states to avoid overfitting.
[States] + ϵ where $\epsilon \sim N(0, \sigma^2)$.
- We don't need to take an action from the agent's q-network every time.
***if random.random* < ϵ_2 :**
take step from the q-network
else:
continue previous action
- Updated ϵ -Greedy policy:
***if random.random* < ϵ_3 :**
take a random step and continue it for n seconds.
else:
take a step from the q-network
- This also helps in reducing the stuttering motion of the agent, make the movement more human-like and helps it getting out of oscillations.
- Noise also helps the agent taking some random actions when it doesn't get any input from the visible screen.

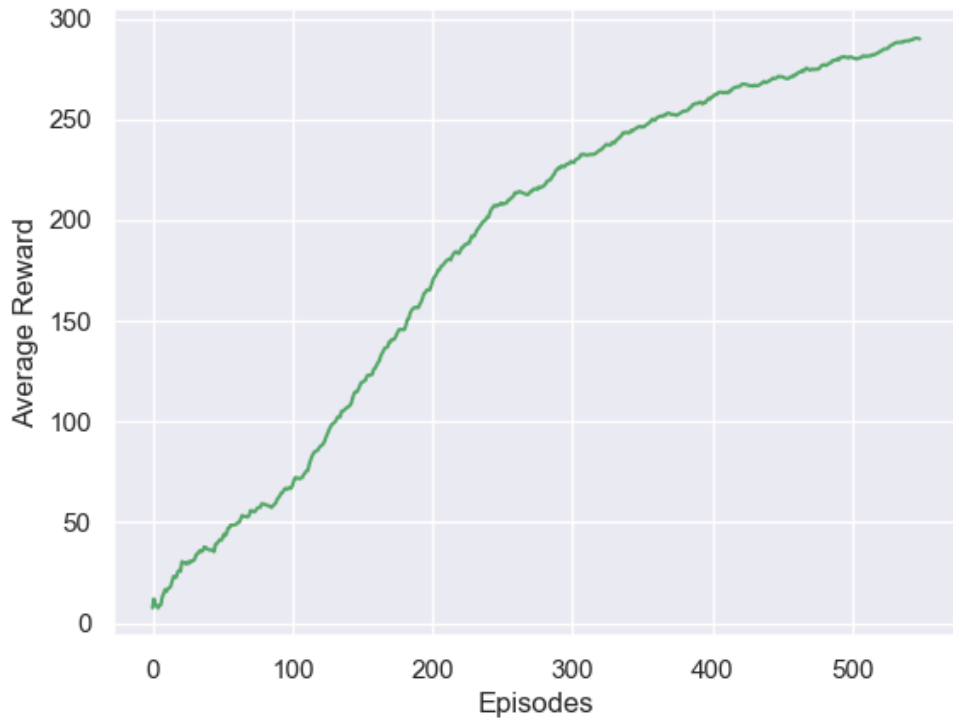
MODEL SUMMARY:

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1, 35, 64]	2,304
Linear-2	[-1, 1, 35, 64]	4,160
Linear-3	[-1, 1, 35, 8]	520
Total params: 6,984		
Trainable params: 6,984		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.04		
Params size (MB): 0.03		
Estimated Total Size (MB): 0.07		

- Size of the model is about 70KB which is very less.
- The model inference time is very fast and doesn't effect or slows down the game.
- This also contributes to agent playing like a human would play.

RESULTS:



Running Average Reward vs Episodes




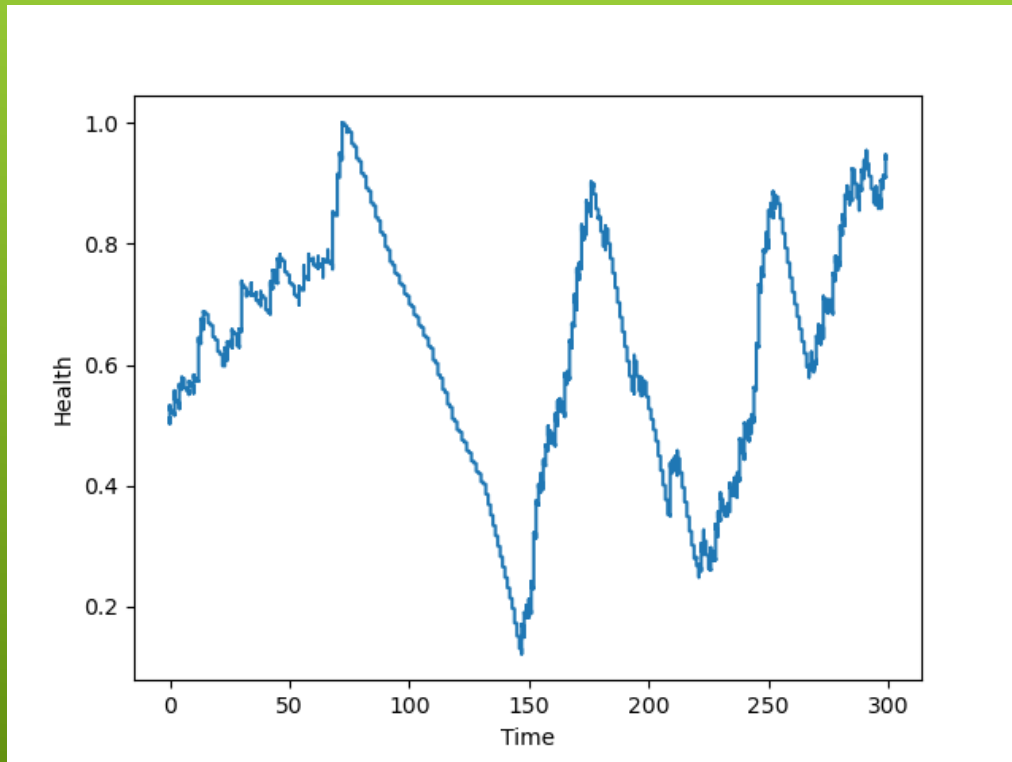
- Time taken for 500 episodes: 26hrs
- The problem's objective starts getting completed after it is able to achieve more than 250 cumulative score.
- The human players average is 344.
- The trained agent average is 416.

Health vs Time:


Same Environment

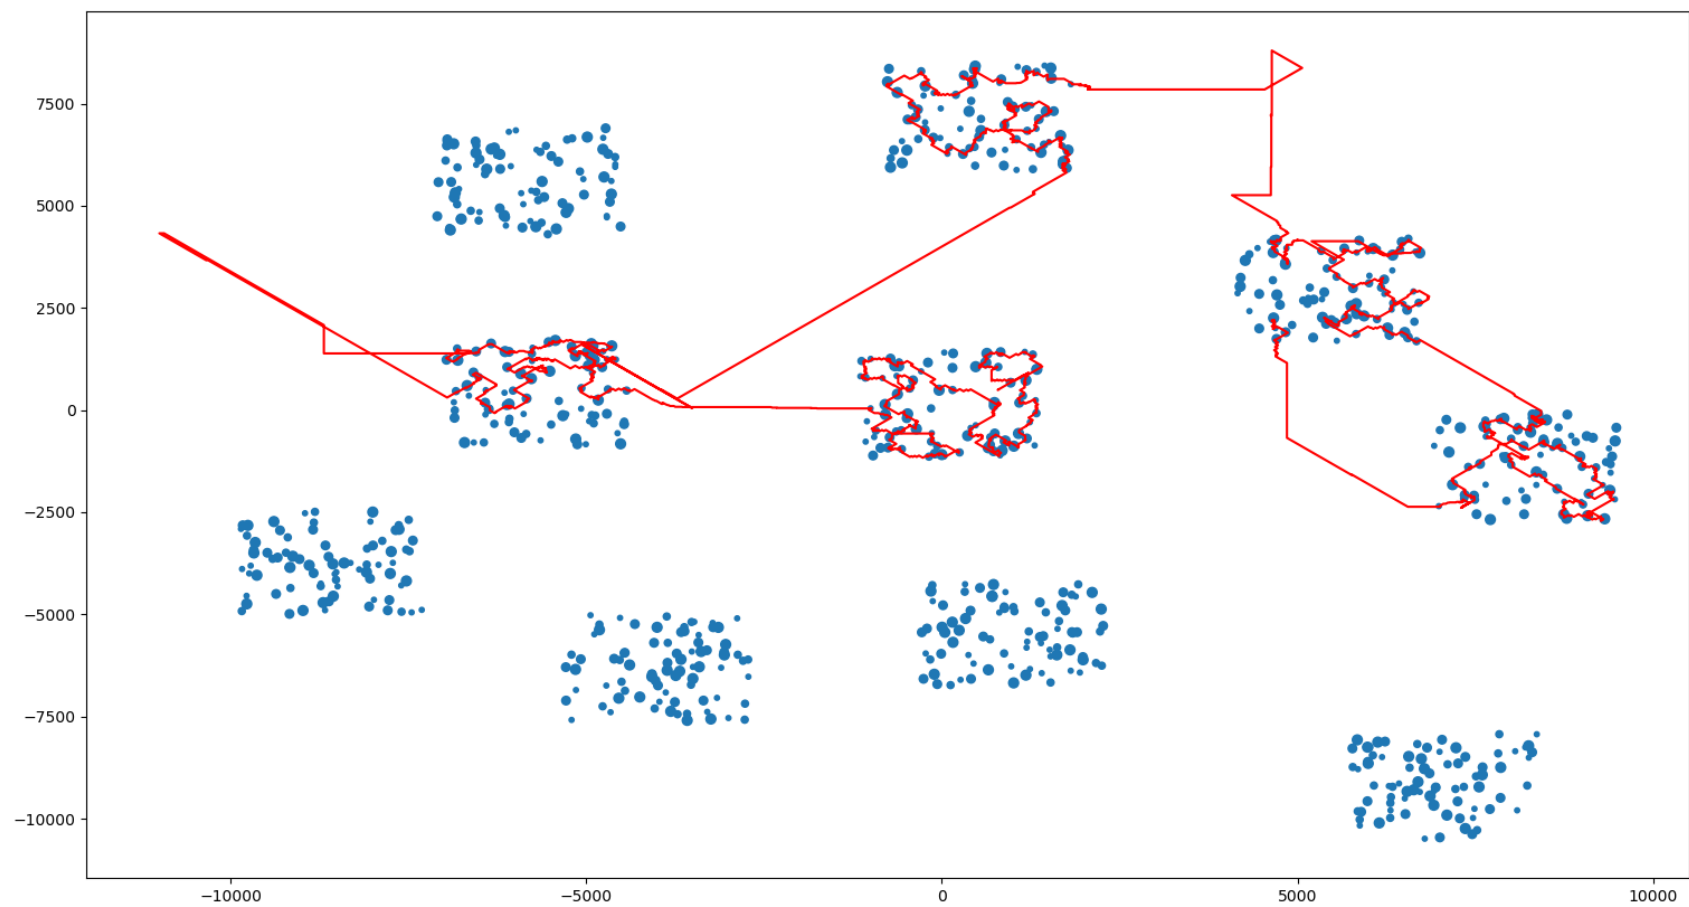


Random Environment

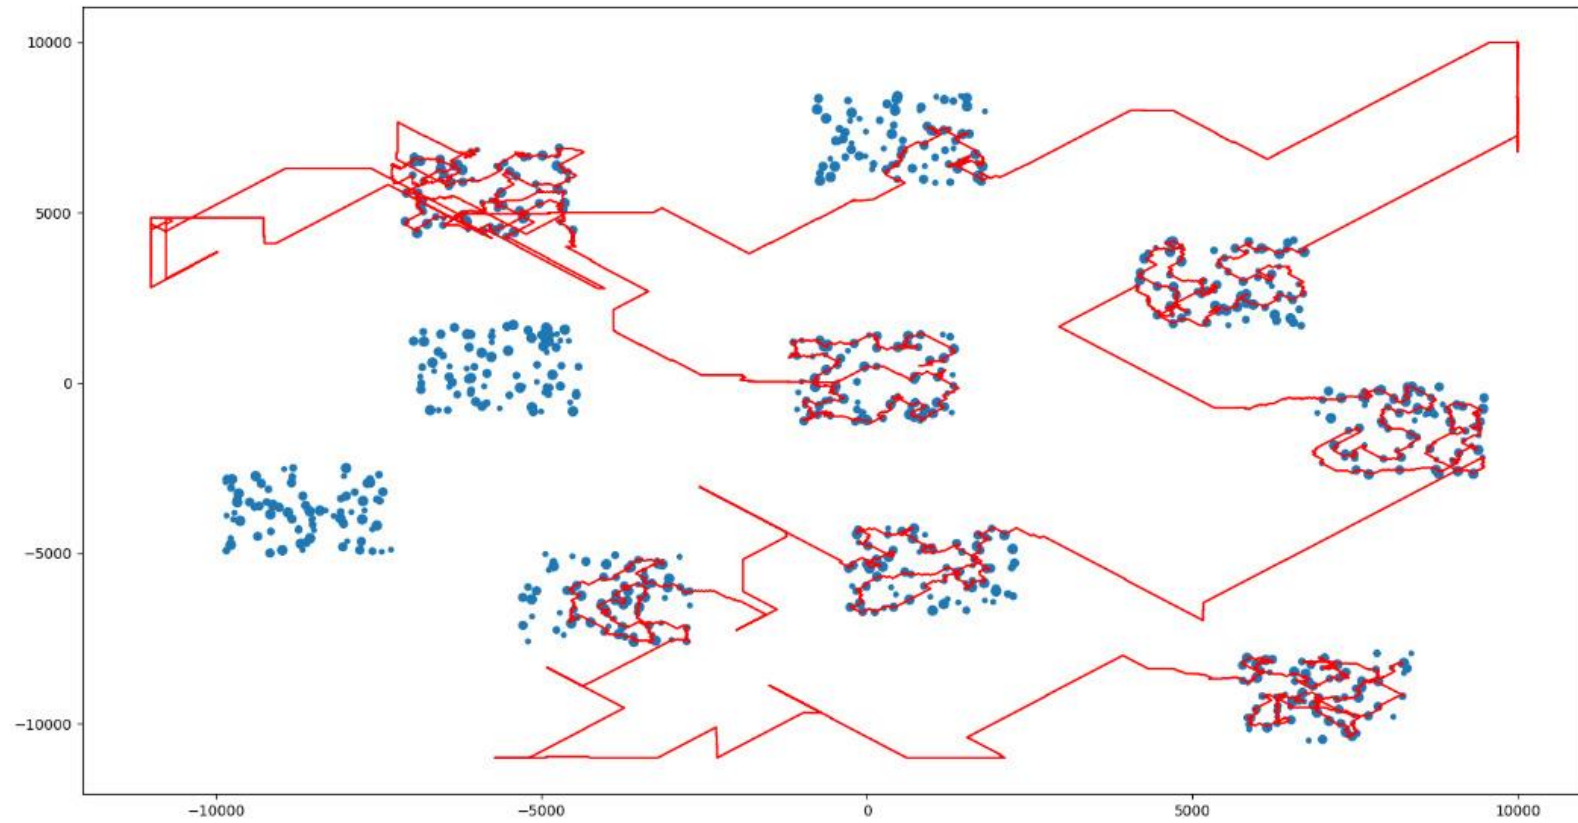


Agent's Path:

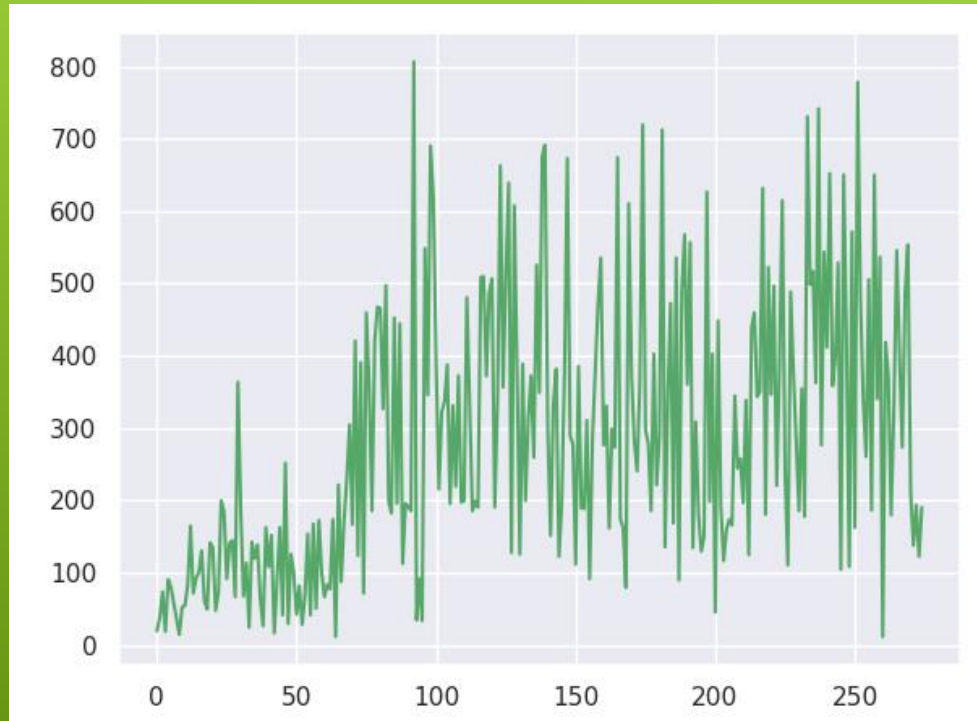




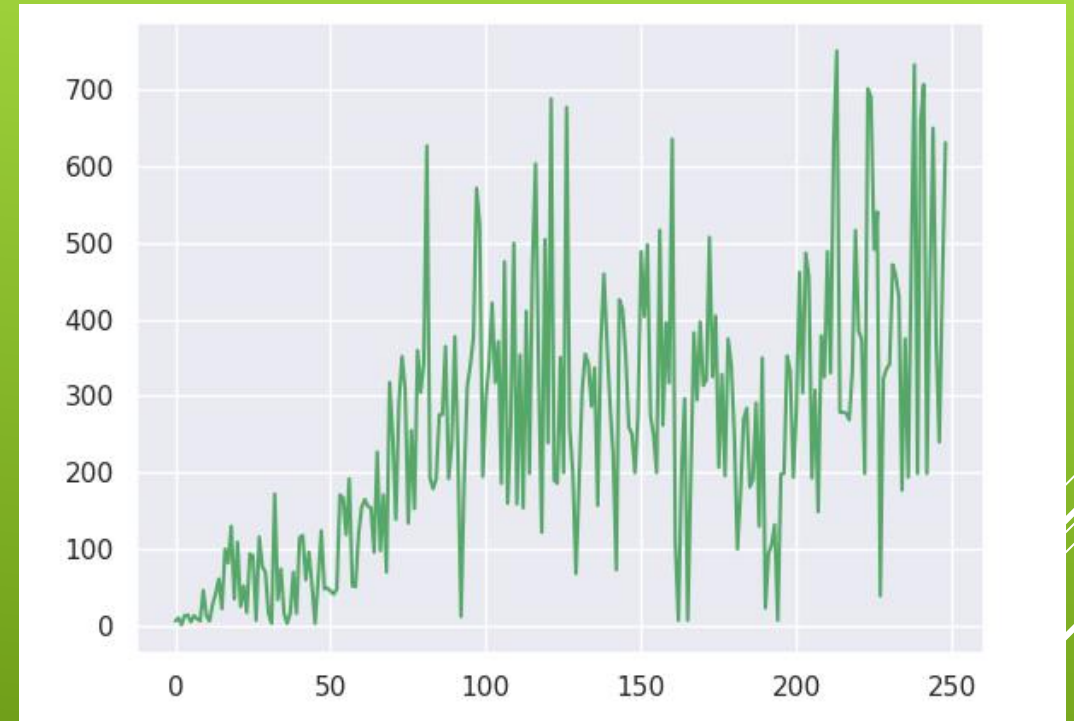
AN OUTLIER !!!!



AGENT'S PERFORMANCE



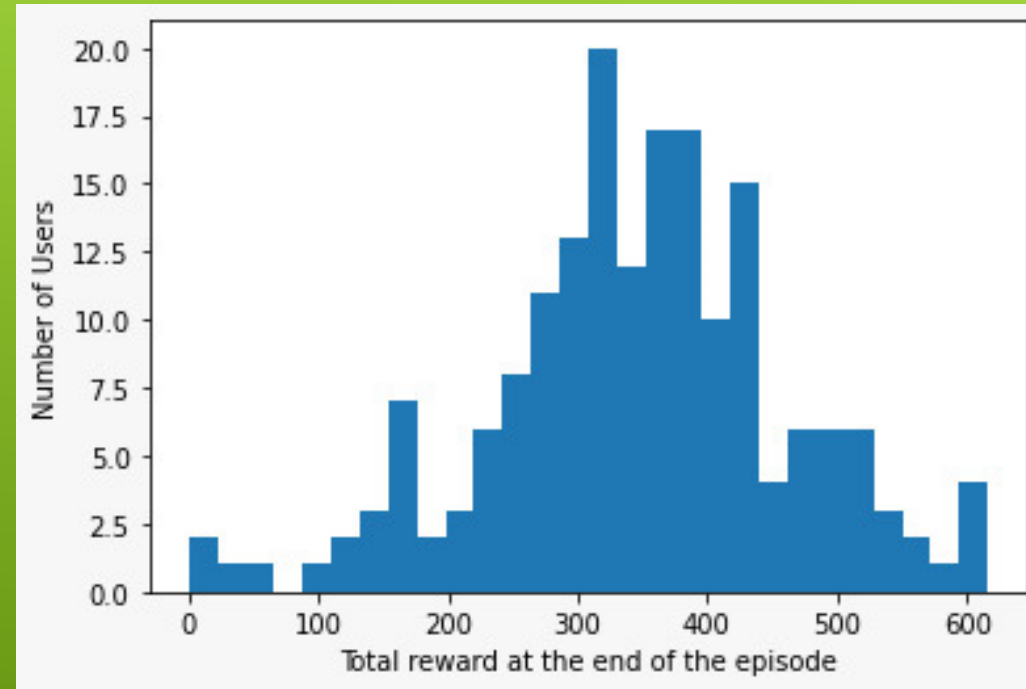
Total rewards vs episodes (Stay option)



Total rewards vs episodes (No-Stay option)

HUMAN PERFORMANCE

The following is the human players performance distribution(data of 183 users.)



CHALLENGES:

- Time to train the agent take about 8-9 hours to reach decent performance.
- As it can be seen, there is lot oscillations ,we believe that it will mitigate as episodes will progress, but it takes lot of time in execution.
- Optimizing the game for smooth performance.
- Resource constraints for using CNN.
- Agents would get stuck into oscillations and the health goes to zero very rapidly.
- Tuning the hyperparameterss of the DDQN.

FUTURE:

- CNN with optimizations of getting the current screen without FPS drop.
 - CNN will help in learning more complicated features.
 - More human-like performance.
- NEAT (Neural Evolution of Augmented Topologies).
 - NEAT is a genetic algorithm: (Simulates the process of natural selection)
 - Individuals with good fitness scores are given preference to pass their genes to successive generations.
 - Cross over of Genes.
 - Mutation in the next generation.
 - Fitness for new population.

FINAL CONTRIBUTION:

NAME	CONTRIBUTION
Tabish Ahmad - 21111060	Environment, Feature selection, Literature survey.
Prakhar Srivastava - 170486	Environment, Pygame and model implementation, Literature Survey, Solution approach, Feature selection, Training agent, getting results, final slides and final report.
Areeb Ahmad -180135	Environment, Solution Approach, Training agent, plotting results, final slides and final report, Data Analysis.
Mohd Niyas P -170394	-
Samarth Mehrotra -20128408	-

DEMO



REFERENCES

- Charnov, E. L. 1976. Optimal foraging: the marginal value theorem. *Theoretical Population Biology* 9:129–136
- Cowie, R. J. (1977) "Optimal foraging in great tits (*Parus Major*)" *Nature* 268:137–139
- Cassini, Marcelo H., Alejandro Kacelnik, and Enrique T. Segura (1990) "The tale of the screaming hairy armadillo, the guinea pig, and the marginal value theorem" *Animal Behavior* 39(6):1030–1050
- A. Mesquita, Y. Nogueira, C. Vidal, J. Cavalcante-Neto and P. Serafim, "Autonomous Foraging with SARSA-based Deep Reinforcement Learning," 2020 22nd Symposium on Virtual and Augmented Reality (SVR), 2020, pp. 425-433, doi: [10.1109/SVR51698.2020.00070](https://doi.org/10.1109/SVR51698.2020.00070).
- https://neat-python.readthedocs.io/en/latest/neat_overview.html

